


Forecasting at Uber: A Brief Survey

Andrea Pasqua

Data Science Manager - Intelligent Decision Systems - Uber



Uber

Time Series at Uber

Use Cases for Forecasting at Uber

- Time series are ubiquitous at Uber


Use Cases for Forecasting at Uber

- Time series are ubiquitous at Uber
 - **Markets:** descriptors for supply, demand, transactions for multi-sided markets




Use Cases for Forecasting at Uber

- Time series are ubiquitous at Uber
 - **Markets:** descriptors for supply, demand, transactions for multi-sided markets



Use Cases for Forecasting at Uber

- Time series are ubiquitous at Uber
 - **Markets:** descriptors for supply, demand, transactions for multi-sided markets




Use Cases for Forecasting at Uber

- Time series are ubiquitous at Uber
 - **Markets**
 - **Internal resources**: compute, storage and data resources

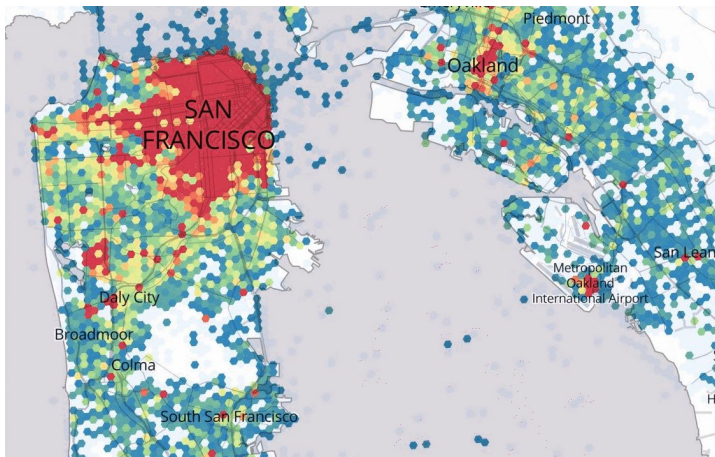
Use Cases for Forecasting at Uber

- Time series are ubiquitous at Uber
 - **Markets**
 - **Internal resources**
 - **Technical time series:** states of micro-services




The Scale of our Domain

- Powerful combinatorics
 - By city, by neighbourhood, by market, by product, by app version
 - By micro-service, by trace, by product, by cohort



<https://eng.uber.com/forecasting-introduction/>




Visualization of a subset of micro-services: edges are API calls

Use Cases for Forecasting at Uber

- Time series are ubiquitous at Uber
- Multiple horizons of interest
 - **Short term:** pricing and reliability
 - **Medium term:** insure markets are well-balanced
 - **Long term:** support Uber's growth without overspending


Compounded Challenges

- High cardinality
- Complex and variable patterns
- Shocks
- Sparse and short time series




Compounded Challenges

- High cardinality
- Complex and variable patterns
- External shocks
- Sparse and short time series



Compounded Challenges

- High cardinality
- Complex and variable patterns
- External shocks
- Sparse and short time series



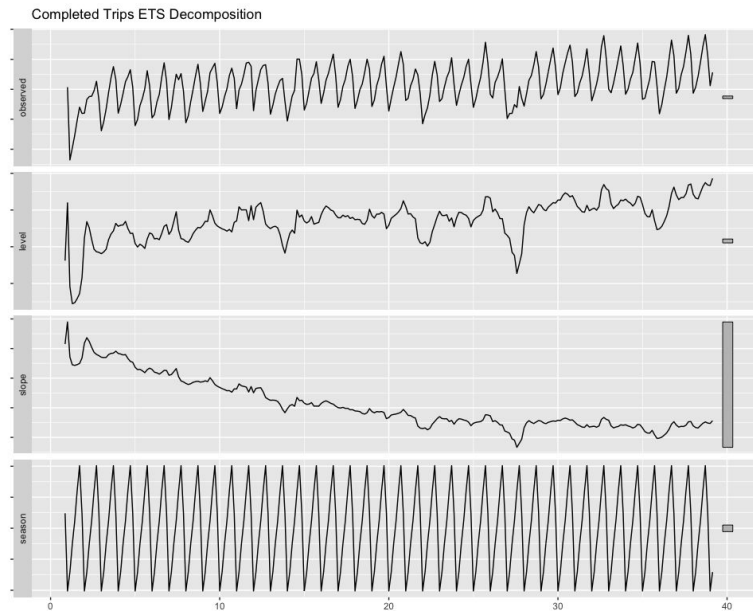
Classical Methods vs. Machine Learning

Classical Methods

- The state space paradigm
 - Decompose the time series into unobservable states
 - Structure equation encoding our beliefs
 - (typically with smoothing)
 - e.g. level, trend, seasonalities
 - Estimate each
 - Recombine and forecast one step at the time

Classical Methods

- The state space paradigm



Classical Methods

- The state space paradigm
- Models in this family are statistical
 - The states can be modeled as stochastic processes
 - Hence you can forecast prediction intervals (a.k.a. density forecasts)
 - A must for informed decision making

Classical Methods

- The state space paradigm
- Models in this family are statistical
- A very successful program!
 - Classical methods are hard to beat when you have
 - Sufficient history
 - Little metainformation
 - Few related time series

A Role for Machine Learning in Forecasting

- But at least one of the conditions above is often violated at Uber
 - Not enough history
 - Short time series in new markets
 - Sparse time series in some markets and for some technical metrics
 - Relevant meta-information
 - Holidays, weather, events
 - Incentives, technological changes
 - Many related time series
 - Uber time series often move with the pulse of the markets

A Survey of ML Techniques

A Survey Based on Personal Experience

- Generalized Linear Model (GLM)
 - The importance of featurization
- Quantile Random Forest (QRF)
 - Algorithms of greater power that learn nonlinearities

A Survey Based on Personal Experience

- Neural Networks (NN)
 - Unexpected challenges
- Hybrid Approaches
 - Best in class

Generalized Linear Model (GLM)


- A linking function to constrain the nonlinear structure

$$y = \lambda^{-1}(w \cdot x)$$


- Stochastic assumptions on the errors
 - Gaussian or Poisson error structure is common
 - Allows density forecasts and optimization via Maximum Likelihood (MLE)

GLM: Application to Incentive Allocations

- Forecasting in the medium term is key to allocate incentives to make our markets healthy and improve customer experience



Forecast (red) vs. Actual (black) for a large cluster



Forecast (red) vs. Actual (black) for a small cluster

GLM: The Importance of Featurization


- This model learns across time series in a given city
- Some of the features convey past observed values of each series ...
- ... but geotemporal information is also a strong predictor
 - Neighborhood
 - Hour of day, day of week, etc.

GLM: The Importance of Featurization

- From one-hot encoding for geotemporal features



- To “fingerprinting”
 - i.e. robust aggregations of observed values by geotemporal units
 - Reduces the complexity of the model, while preserving all the relevant information




GLM: Aggregate Performance

- Incremental improvement over traditional methods
 - wMAPE: weighted Mean Absolute Percent Error is a common error metric

Error Metric	Baseline	GLM
wMAPE	-	↓ 12.5%

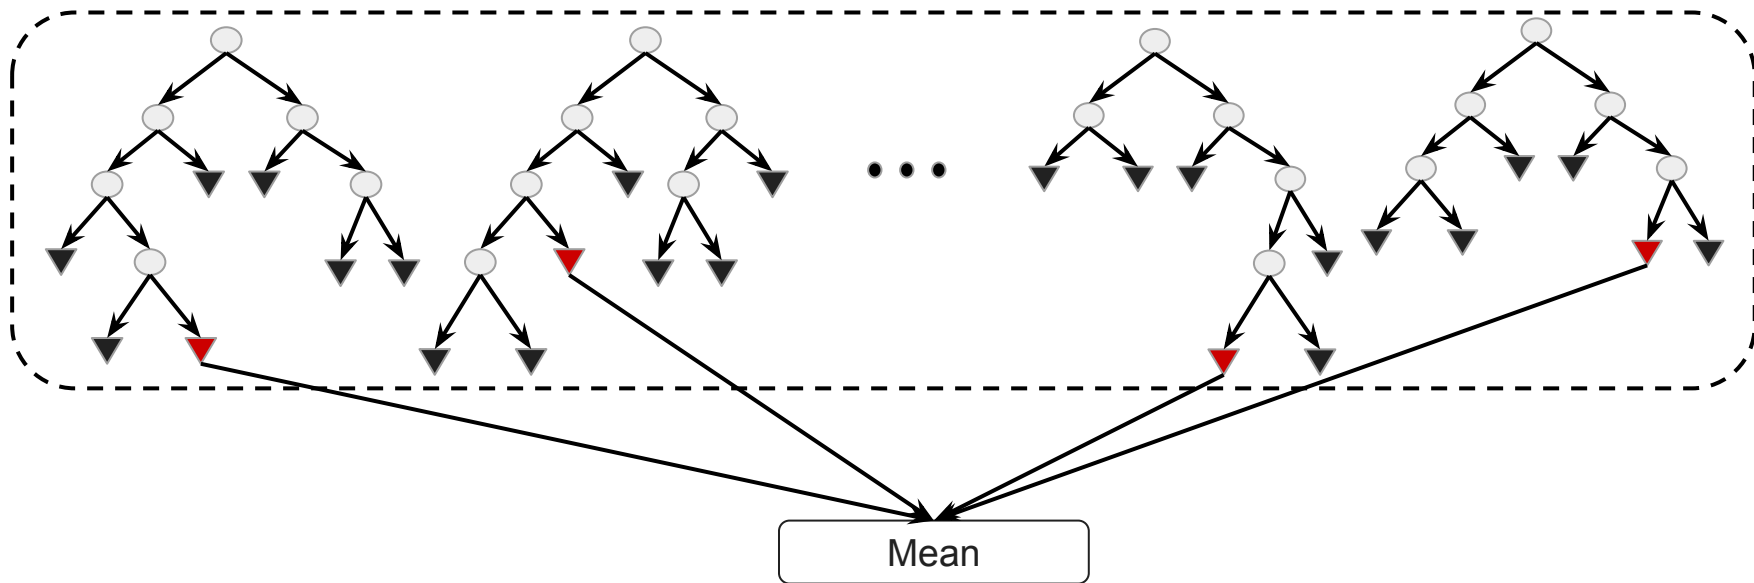
Quantile Random Forest (QRF)

- Decision Trees
 - Greedy optimizer for Purity/Variance




Quantile Random Forest (QRF)

- Add randomness and ensemble them in a forest
 - Less variance



Quantile Random Forest (QRF)

- From point to density forecast
 - Essential for informed decision-making



QRF: some remarks

- Nonlinearities can be learned (ML proper)
- Featurization is again paramount and it is hard work...
 - No sense of time in the algorithm itself
 - Yet easier than for Neural networks
 - Forgiving to categorical and feature with broadly different codomains


QRF: Aggregate performance

- Proved remarkably effective!

Error Metric	Baseline	GLM	QRF
wMAPE	-	↓ 12.5%	↓ 22%


Neural Networks (NN)

- Architecture of interconnected neurons with nonlinear activation functions and weights that can be learned
- A variant, Recurrent Neural Networks (RNN), has a sense of time!
 - An internal mutable state that allows for memory



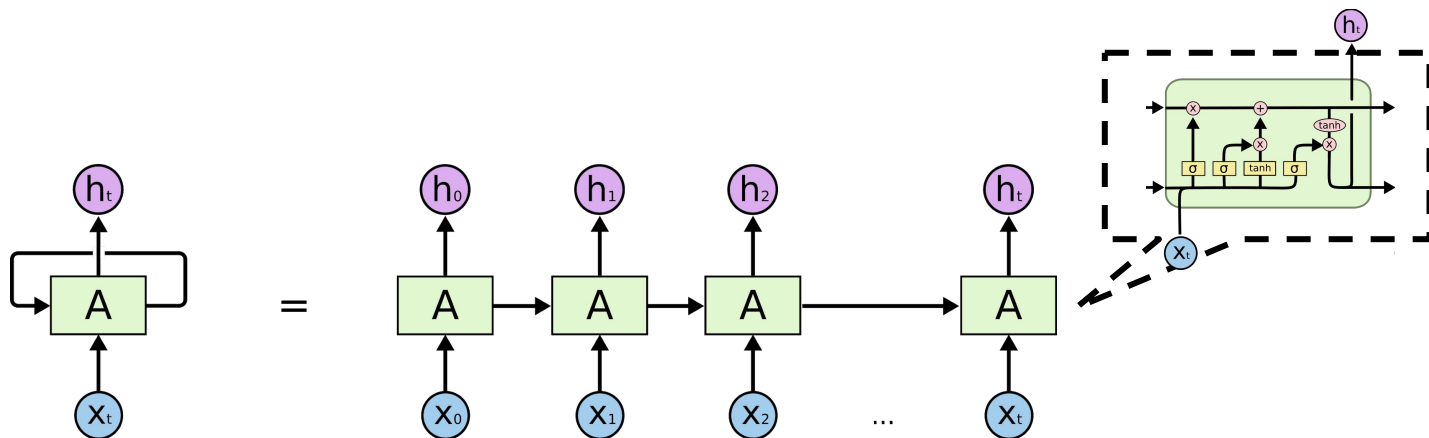
Long Short-Term Memory (LSTM)

- Long-, Short-Term Memory cells are often used to build RNNs
 - Retain variable amount of past information
 - Fight off vanishing gradients




Long Short-Term Memory (LSTM)

- LSTM Cells have an internal structure



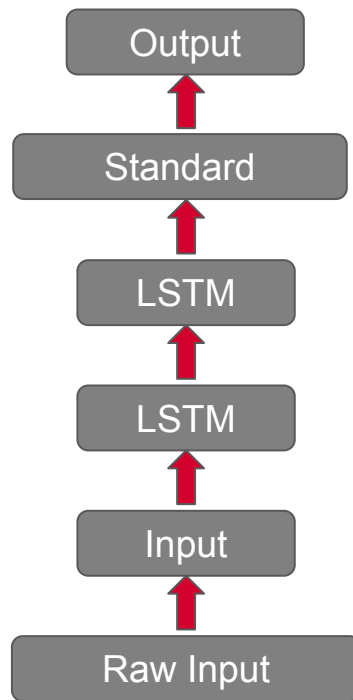
Long Short-Term Memory (LSTM)

- An LSTM cell has internal structure
 - Three gates
 - Forget
 - Input
 - Output
 - Allows for selective and adaptive memory



NN: Typical Architectures

- Not very deep architecture
 - If compared with computer vision
 - Many variations for different use cases
- Preprocessing is required and a challenge
 - Normalizing time series and features
- Used for
 - Incentive allocations
 - Anomaly detection
 - As part of hybrid models




NN: Aggregate Performance

- Robust performance, but not an improvement over QRF
- Is this the final word?

Error Metric	Baseline	GLM	QRF	MLP	LSTM
wMAPE	-	↓ 12.5%	↓ 22%	↓ 12.5%	↓ 18.8%

NN: Recovering Density Forecasts

- Optimizing with a pin-ball loss estimates quantiles
- Monte Carlo drop-out method
 - [Engineering Uncertainty Estimation in Neural Networks for Time Series Prediction at Uber -- Uber Eng Blog](#)



Hybrid Models

Hybrid Models

- They combine
 - Classical structural models (e.g. state space models)
 - Machine Learning models (e.g. Deep learning)
- They are hierarchical
 - Some of the parameters are local, i.e. fit per time series...
 - ... others are global, i.e. learned across time series
 - The local parameters provide the specificity NNs often lack

An Event Lift Model

- Uber time series are significantly impacted by external events
 - Holidays
 - Extreme weather
- Classical models do well on ordinary days but have trouble learning from rare events

An Event Lift Model

- A hybrid solution
- A Generalized Additive Model (GAM) for the baseline forecast
 - Six months out
 - Detrending with exponential smoothing
 - Capturing seasonality with semi-parametric splines
- Quantile Random Forest to learn multiplicative event lifts across time series


An Event Lift Model

- Improvement of six months ahead forecast for number of trips

Error Metric	Region	HW	GAM	GAM + QRF
wMAPE	US	-	↓ 8.1%	↓ 13.5%
wMAPE	Latin America	-	↓ 32.2%	↓ 40.4%


The M4 model

- The M4 model: an ES-RNN hybrid




The M4 model

- The M4 model: an ES-RNN hybrid




The M4 model

- The M4 model: an ES-RNN hybrid




The M4 model

- The M4 model: an ES-RNN hybrid



The M4 model

- The M4 model: an ES-RNN hybrid




- Exponential smoothed structural components
- No need to capture local trend


Dynamic Computational Graphs

- Two types of parameters
 - The NNs weights \mathbf{W} are learned across time series (robustness)
 - The smoothing coefficients α and γ time series by time series (specificity)
- The computational graph will vary from time series to time series
 - This mode is supported by several frameworks: Dynet, PyTorch and more recently TensorFlow

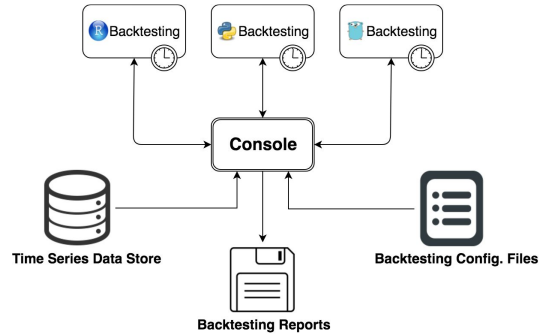
Learn more about Forecasting at Uber



[M4 Forecasting Competition: Introducing a New Hybrid ES-RNN Model](#)



[Engineering Uncertainty Estimation in Neural Networks for Time Series Prediction at Uber](#)



[Omphalos. Uber's Parallel and Language-Extensible Time Series Backtesting Tool](#)



[Forecasting at Uber: An Introduction](#)

Thank you!



Proprietary and confidential © 2018 Uber Technologies, Inc. All rights reserved. No part of this document may be reproduced or utilized in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval systems, without permission in writing from Uber. This document is intended only for the use of the individual or entity to whom it is addressed and contains information that is privileged, confidential or otherwise exempt from disclosure under applicable law. All recipients of this document are notified that the information contained herein includes proprietary and confidential information of Uber, and recipient may not make use of, disseminate, or in any way disclose this document or any of the enclosed information to any person other than employees of addressee to the extent necessary for consultations with authorized personnel of Uber.